

Reconciling Eventually-Consistent Data with CRDTs

Starring **Noel Welsh**

A **myna** Production

In conjunction with
_underscore

Showing at

Scala eXchange 2013

A detail from Michelangelo's 'The Last Judgment' fresco in the Sistine Chapel. It depicts a scene of divine judgment where the dead are being resurrected and condemned. In the center, Christ sits enthroned, surrounded by angels. Below him, figures are shown in various states of torment, some being pulled apart by demons. The scene is filled with dramatic lighting and anatomical detail.

**Why are we
here?**

CRDTs

(An overview)

Conflict-free Replicated Data Type

Conflict-free Replicated Data Type

Conflict-free Replicated Data Type

Conflict-free Replicated Data Type

**Conflict-free
Convergent
Commutative**

Merge data
automatically

I

Scala

#2

Abstract Algebra

#3

Special Relativity

**Why do we
care?**

You have an
awesome
web site

So you want
sub-Second
page load

Use Scala

Spray: >200K requests/s

Rails: 4640 requests/s

Source: Tech Empower JSON serialization benchmark

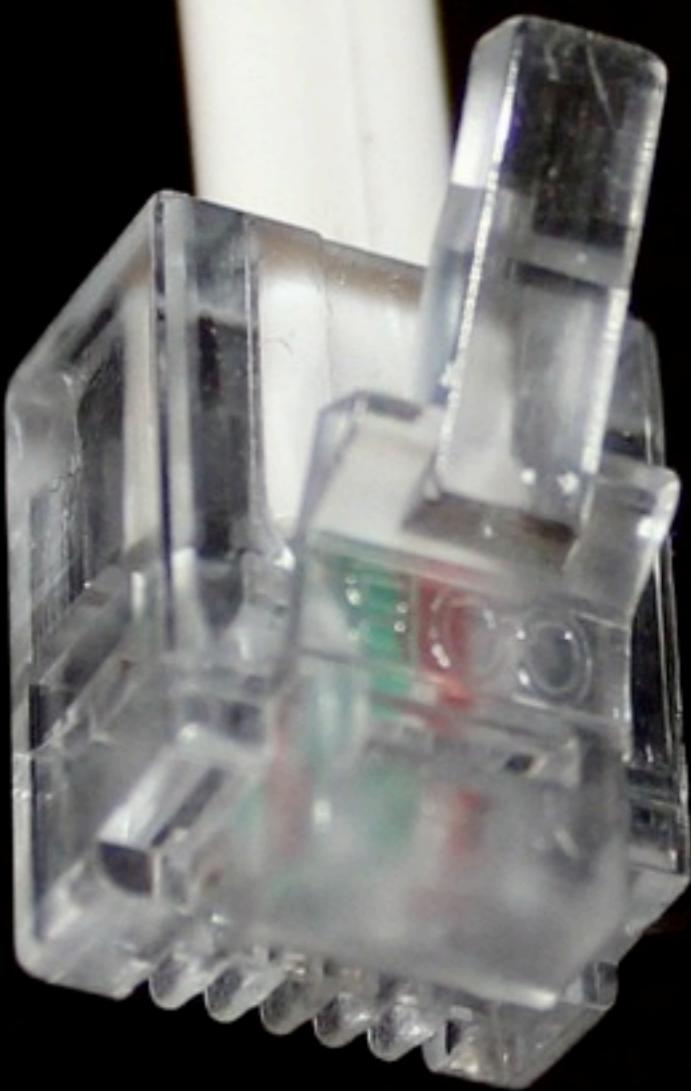
<http://www.techempower.com/benchmarks/>

That's
not
enough

E=MC²



World map from Wikimedia Commons

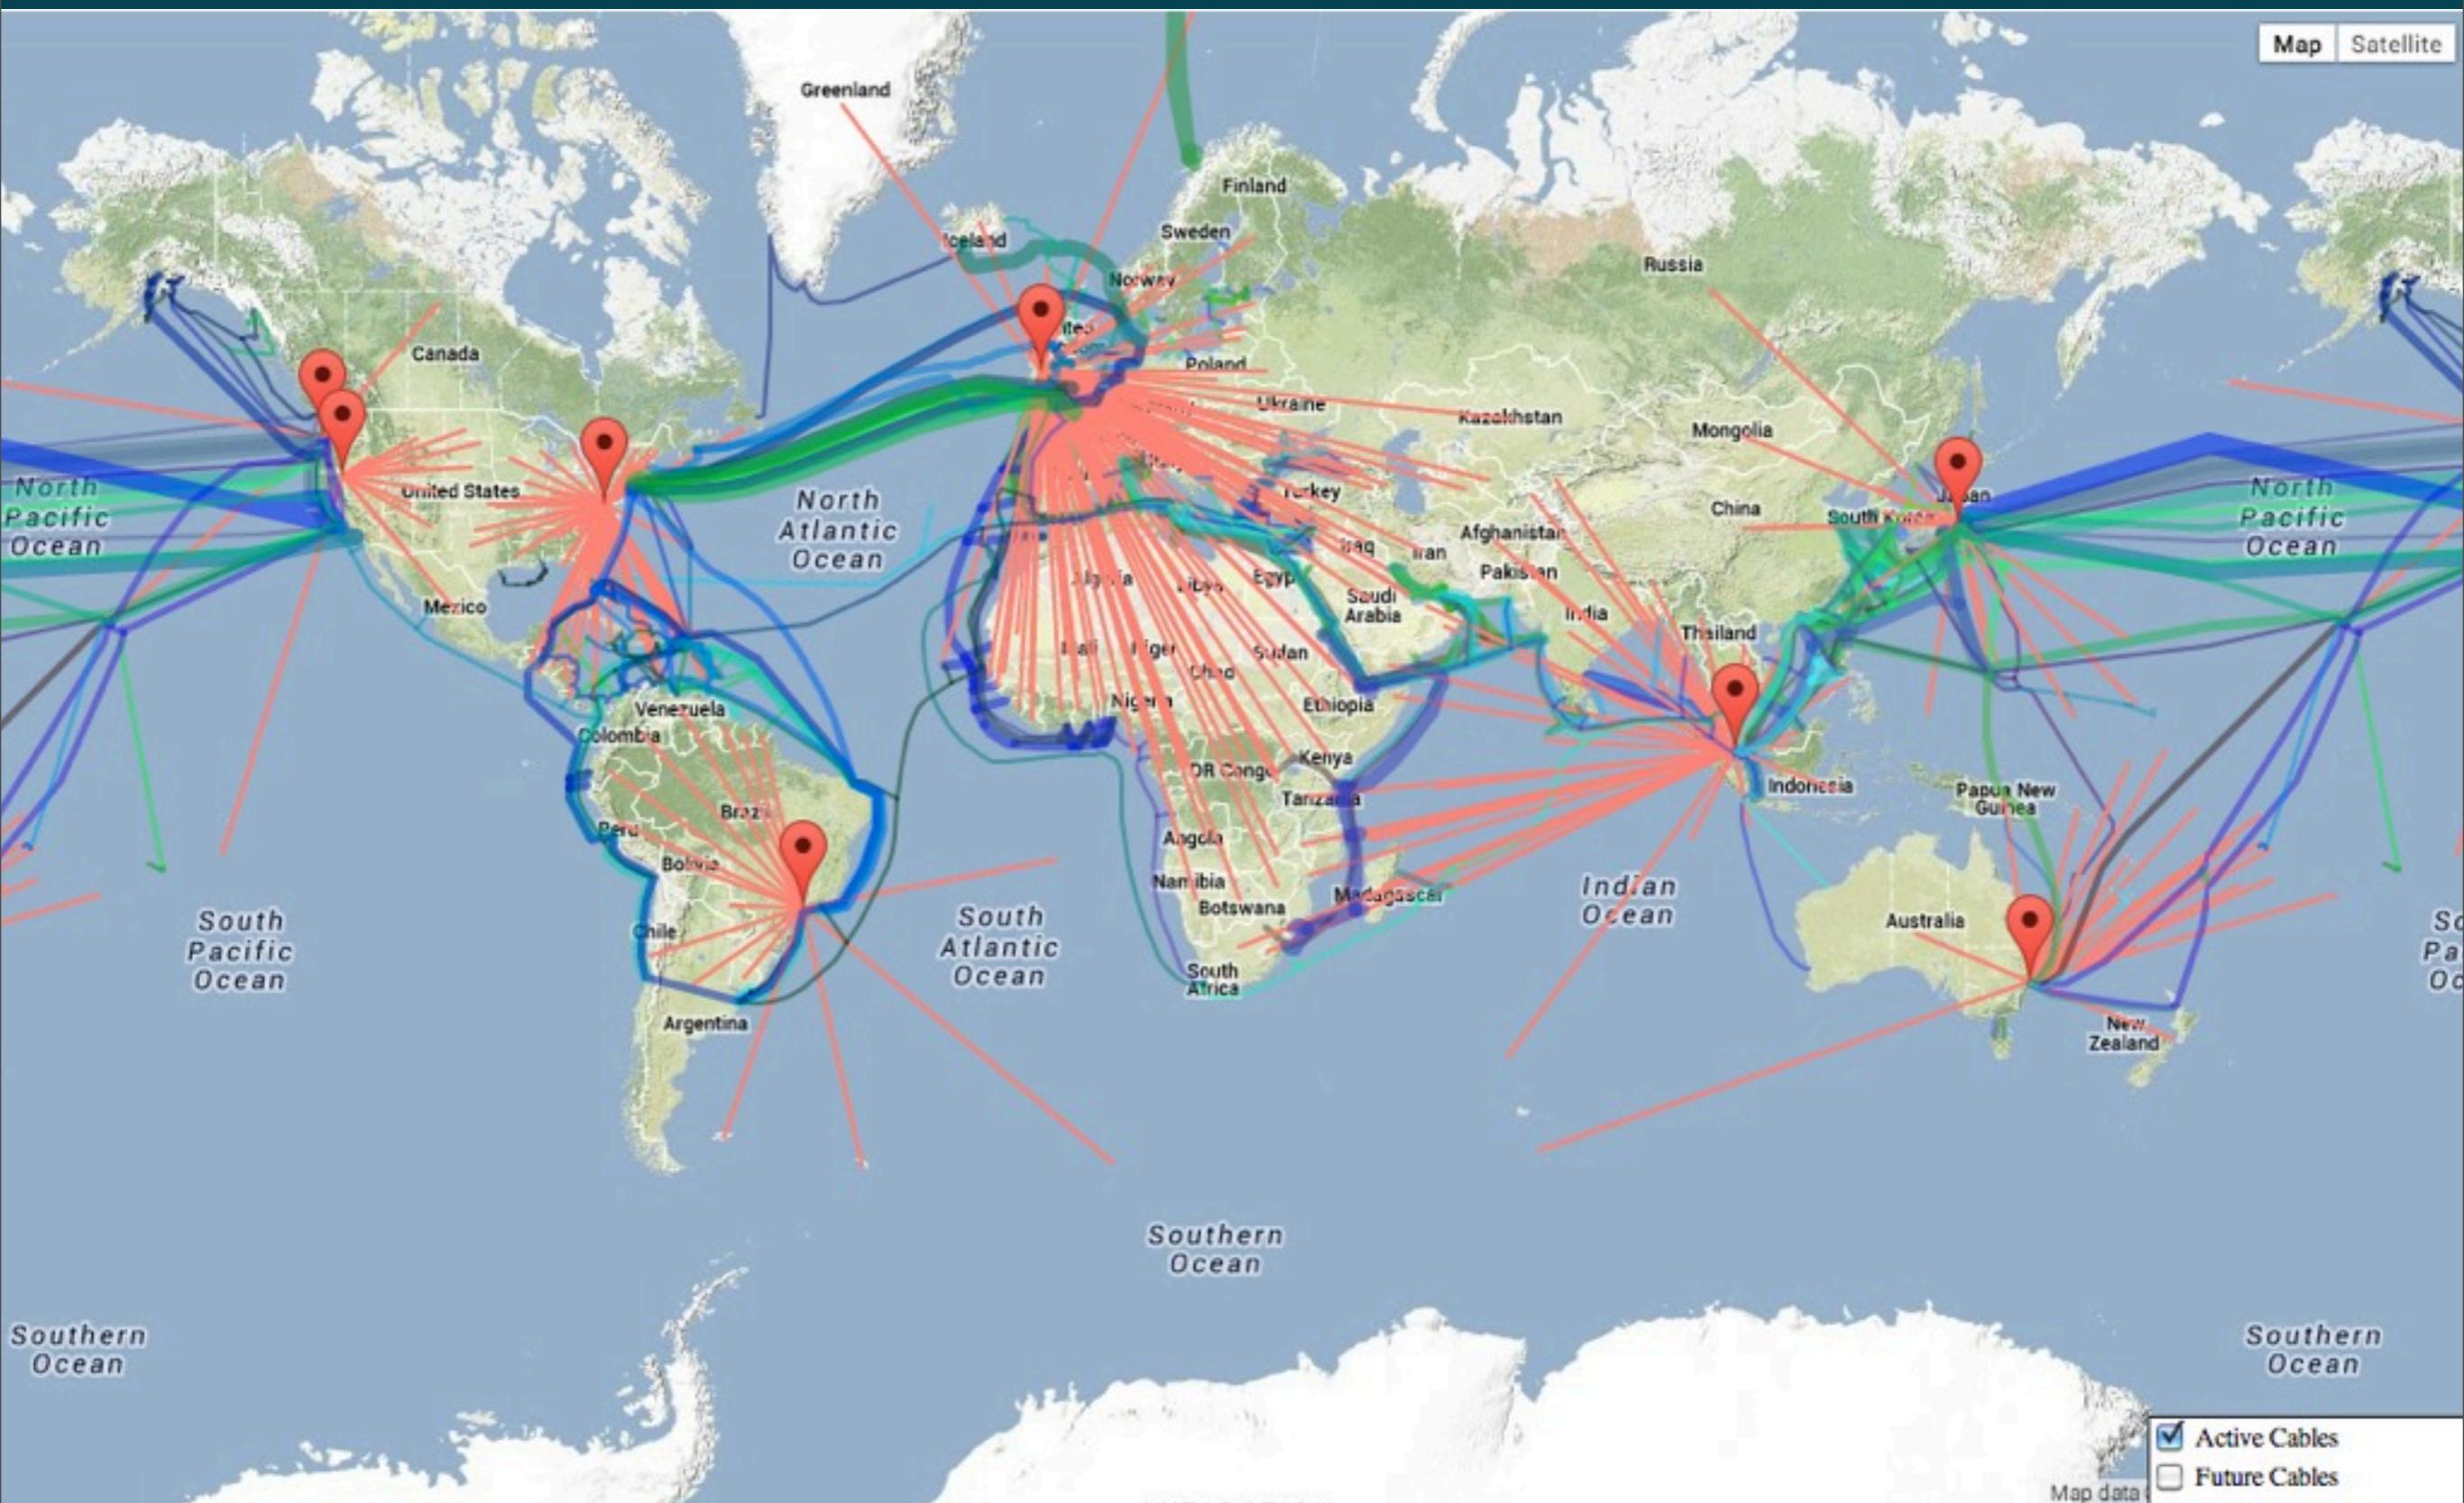


<http://www.flickr.com/photos/21561428@N03/5185781936/>

LOCATION

LOCATION

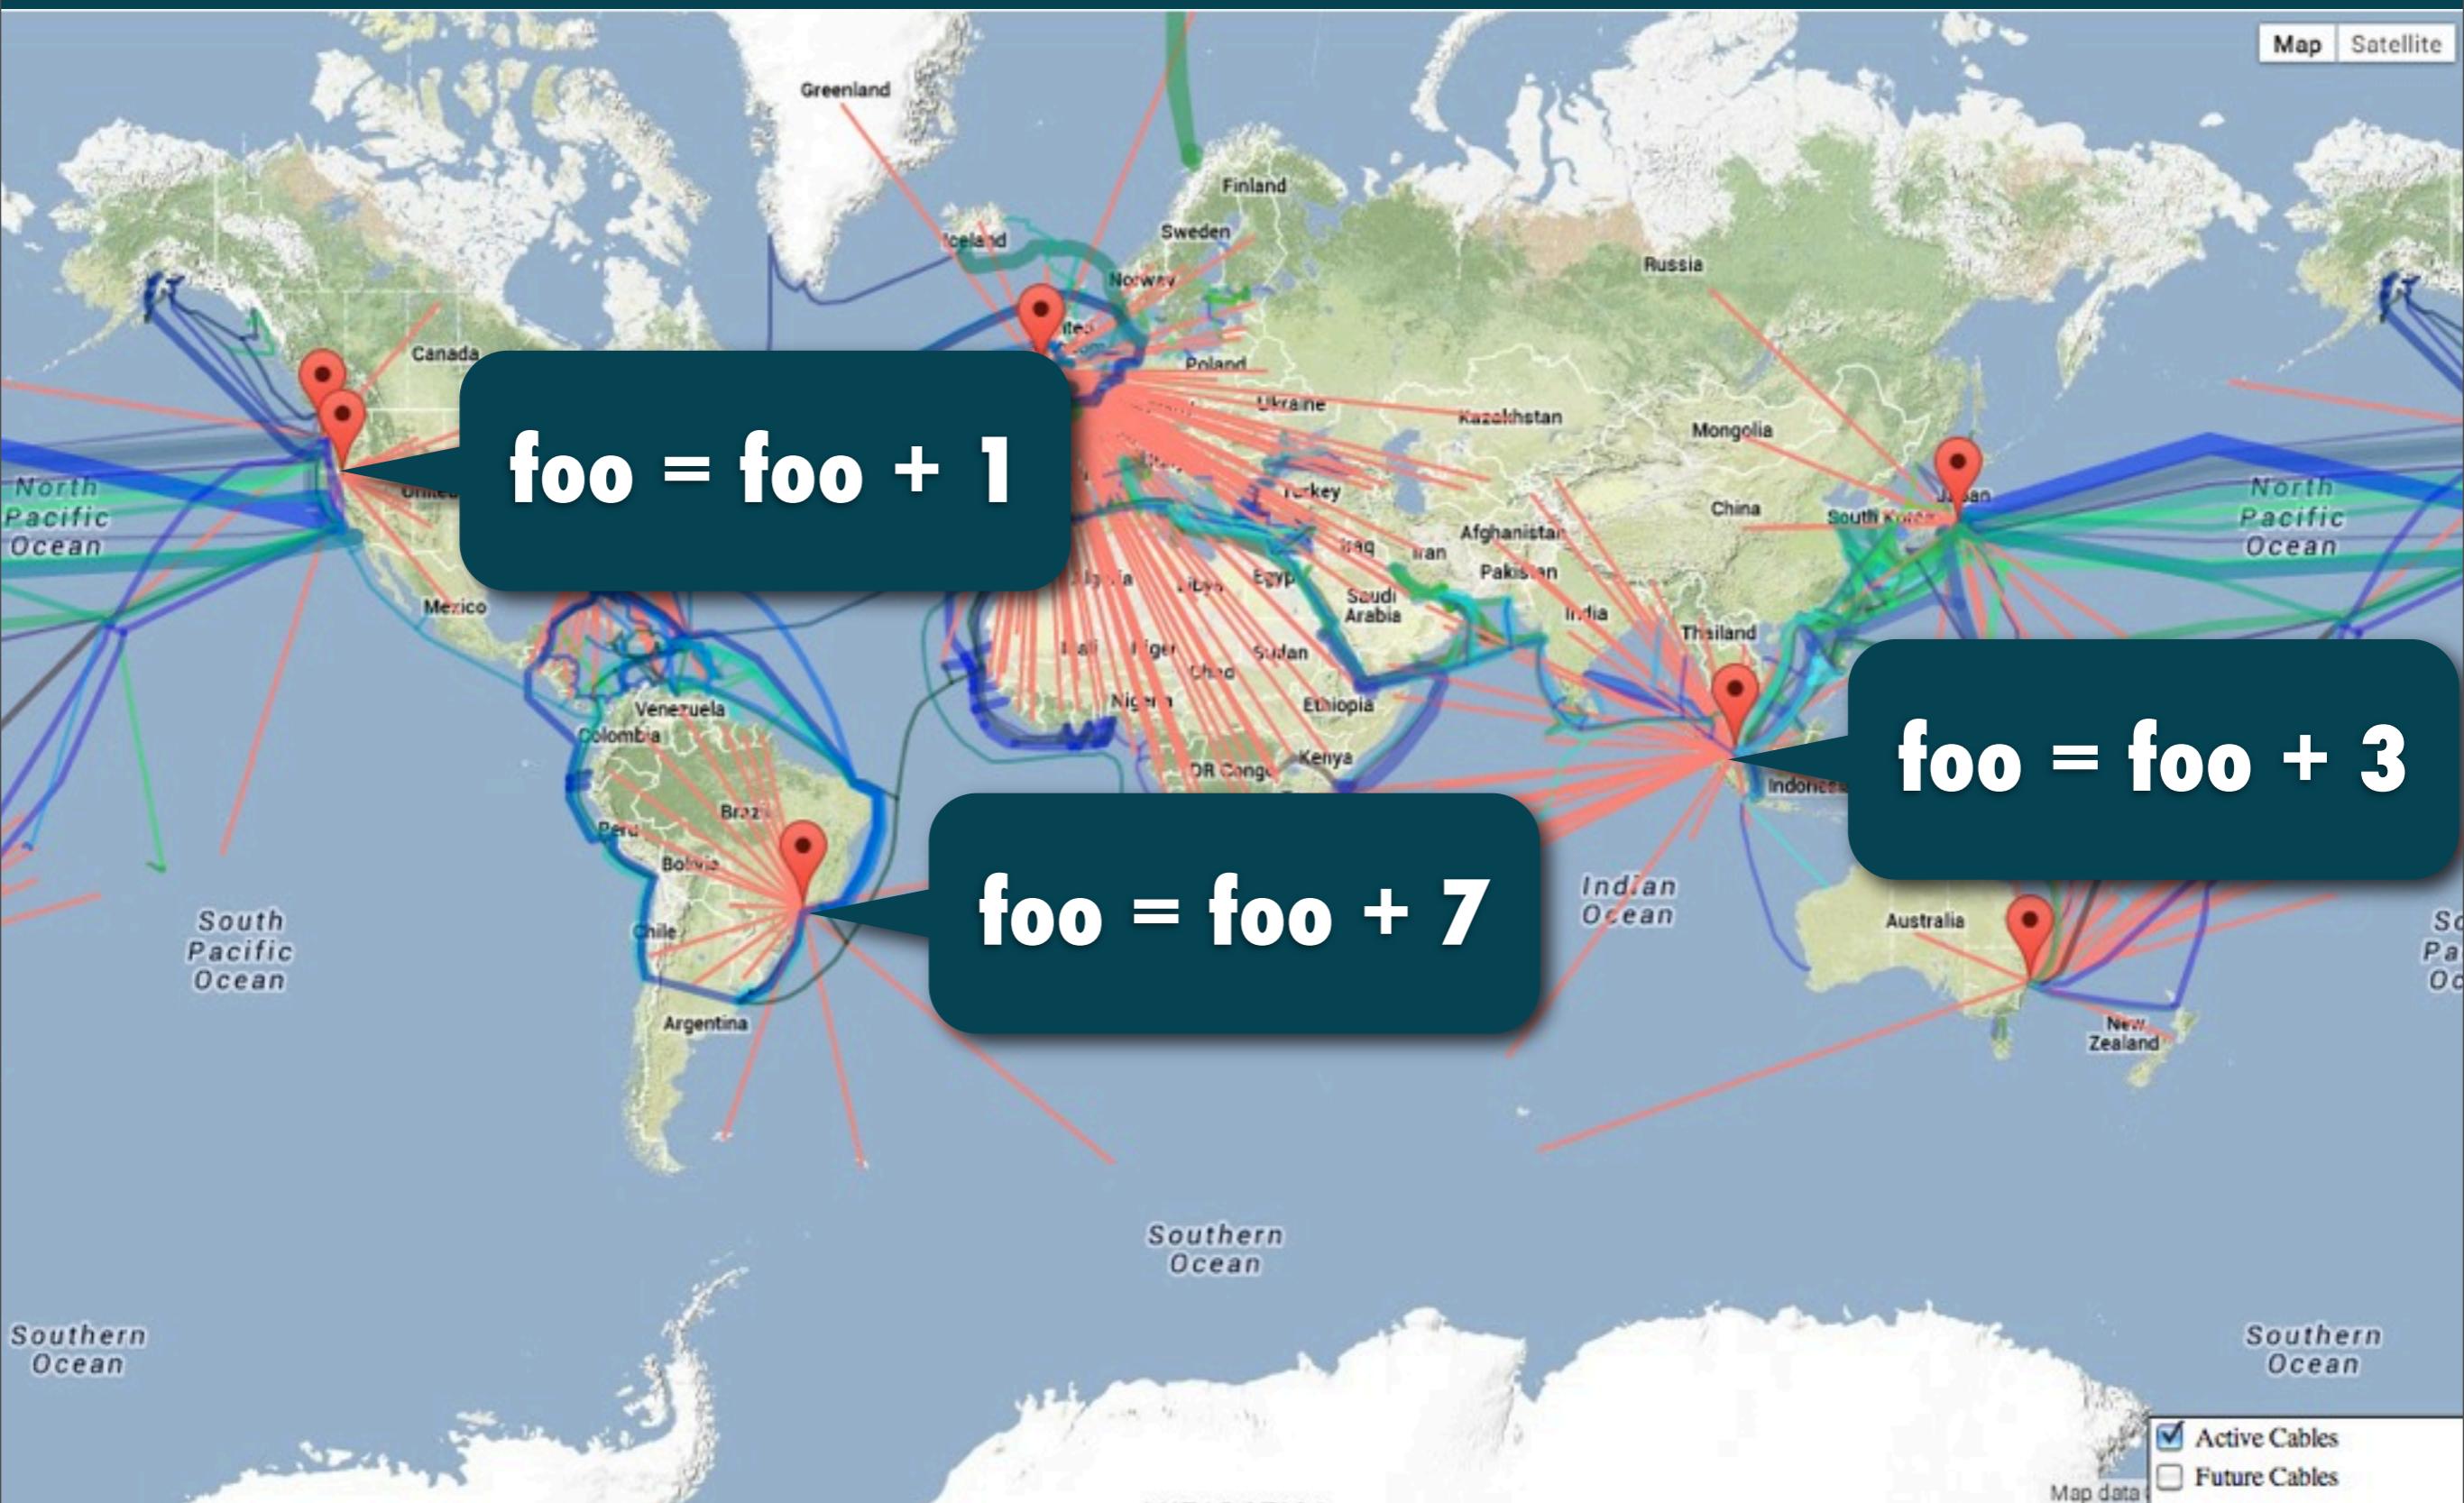
LOCATION



<http://turnkeylinux.github.io/aws-datacenters/>

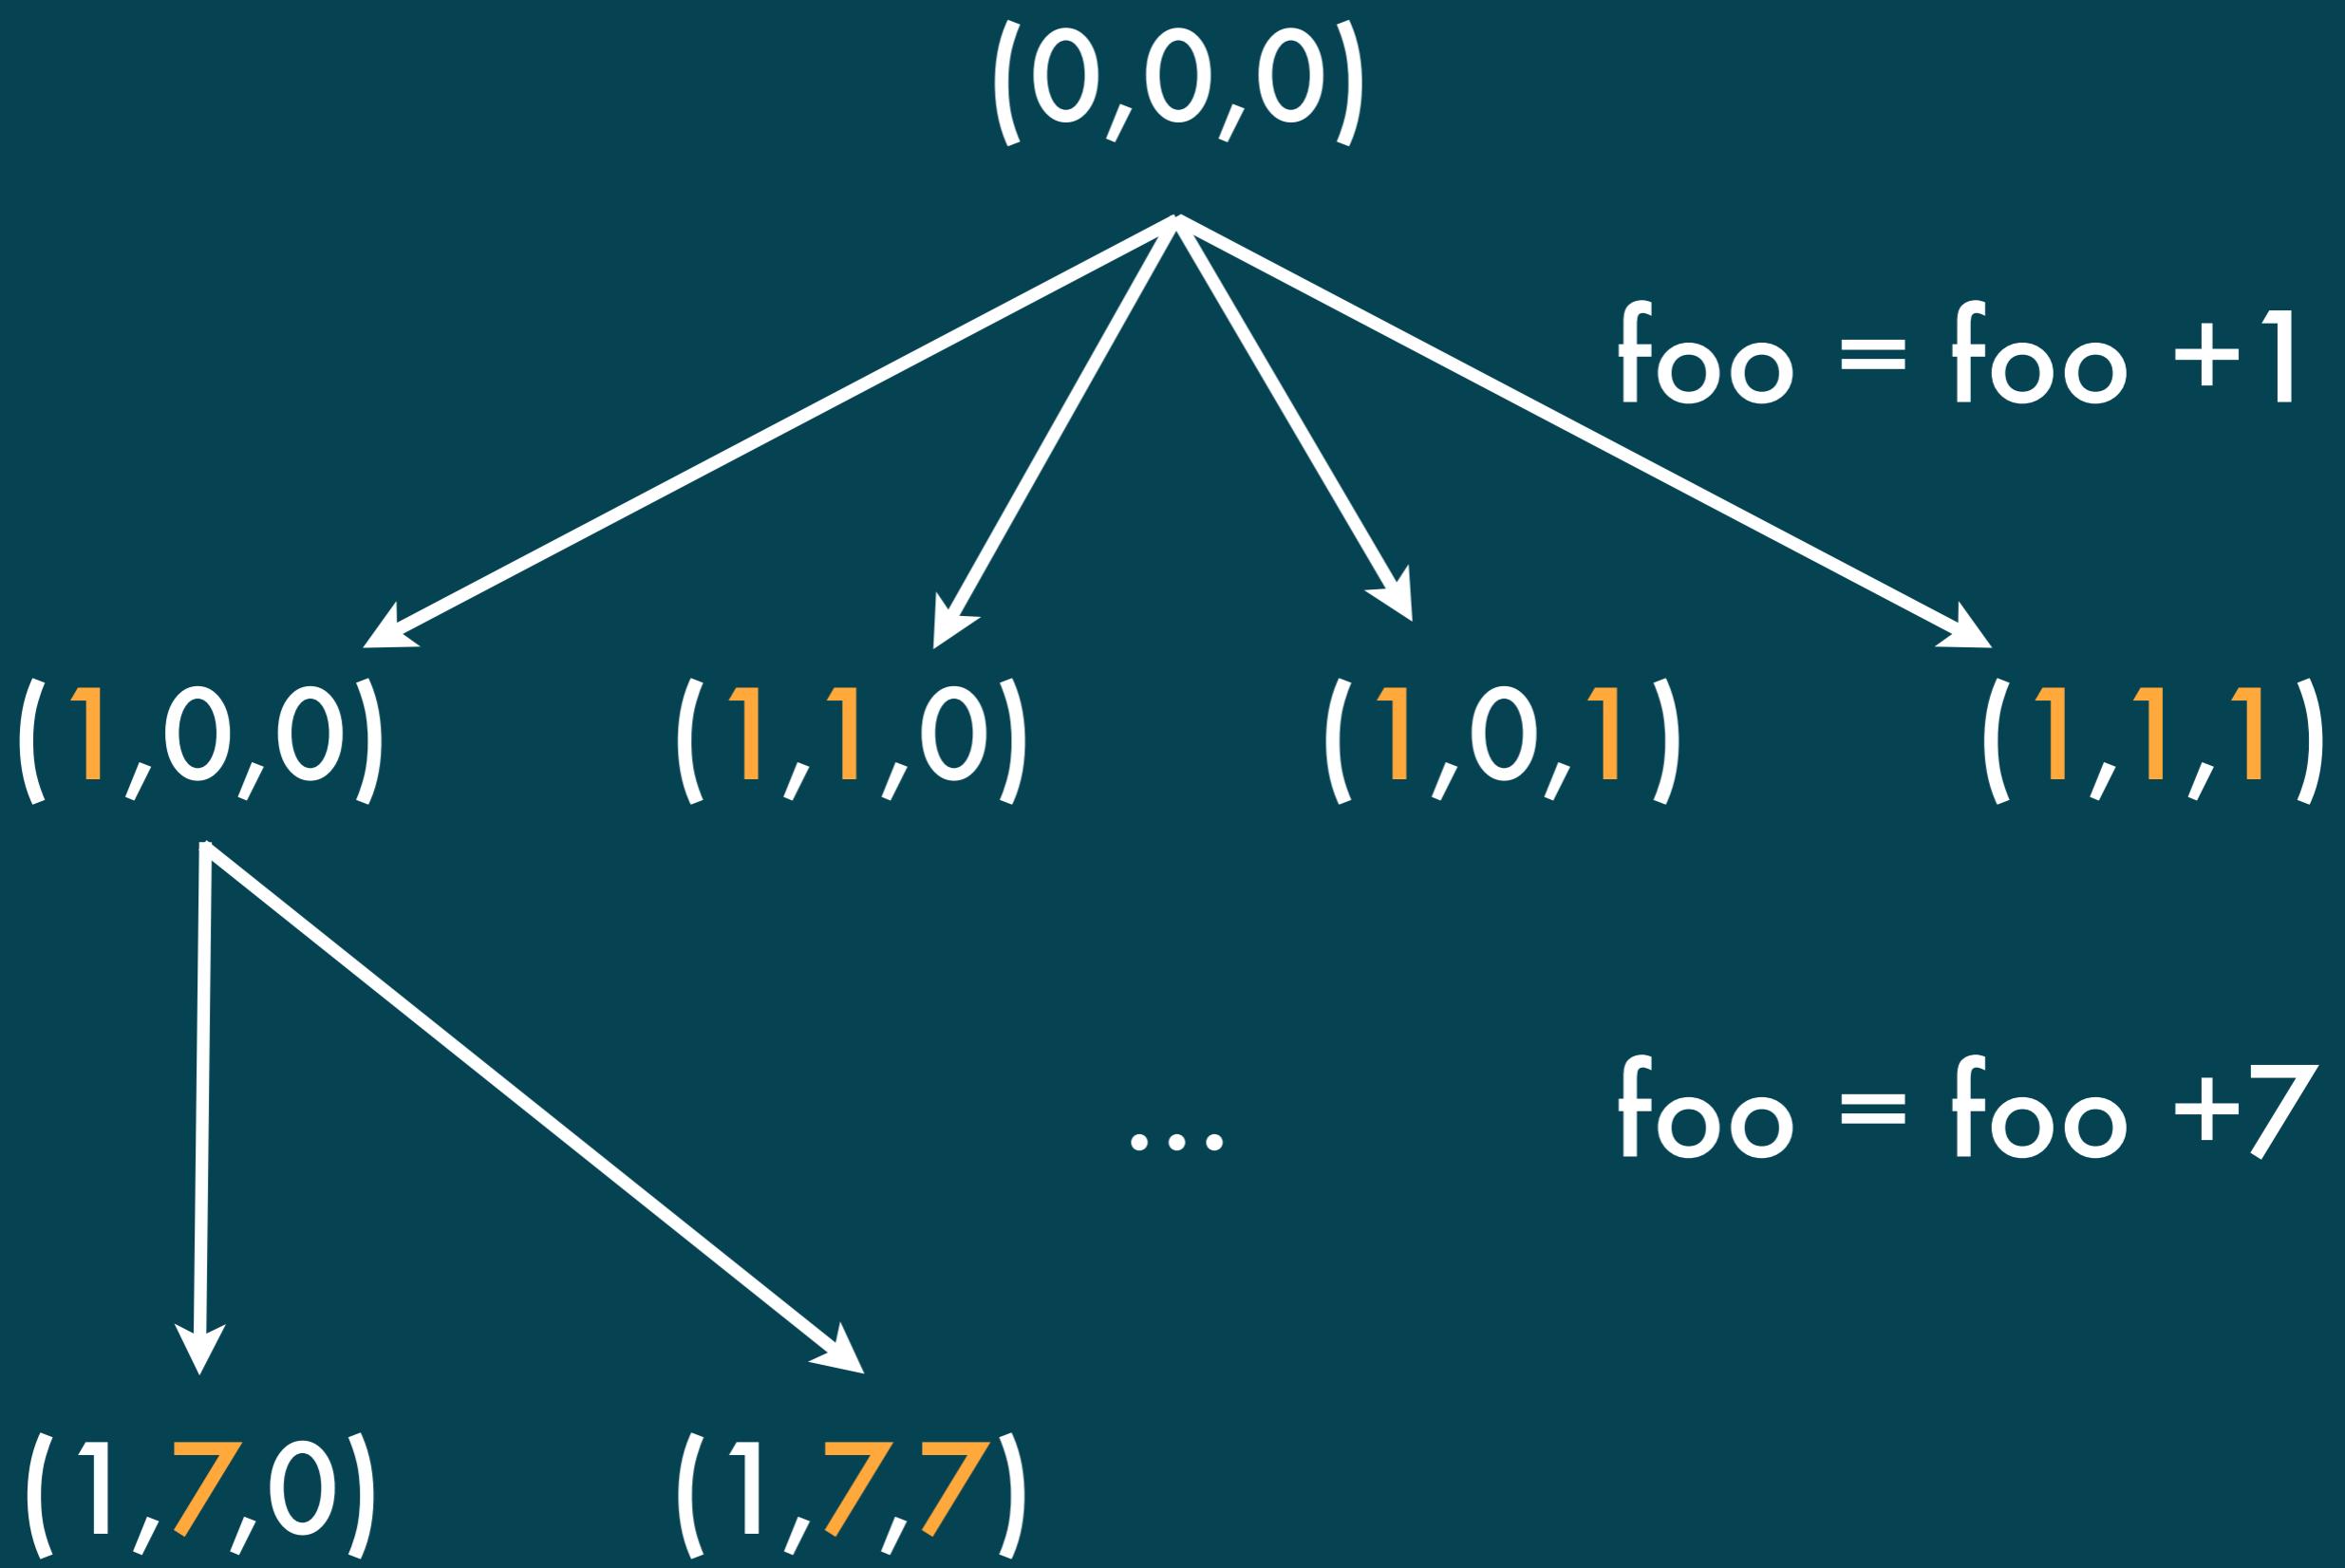
**Problem
SOLVED!**

**Problem
SOLVED?**



<http://turnkeylinux.github.io/aws-datacenters/>

What the FOO?



We have conquered
Latency
We have lost
Consistency

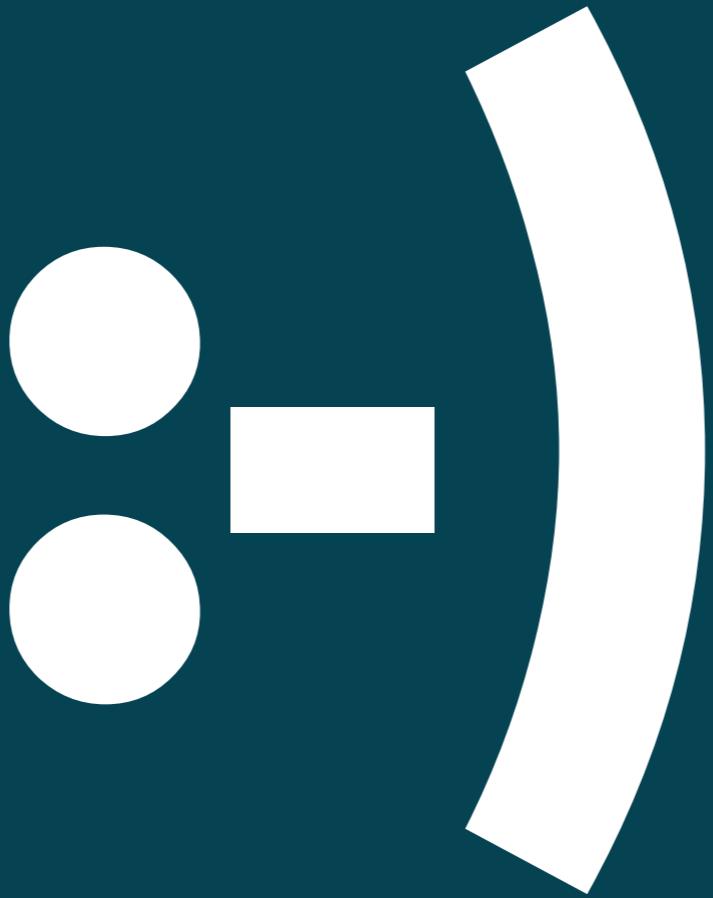
Solution #1 (Quantum Mechanics)

Simply consider all possible world states and ...



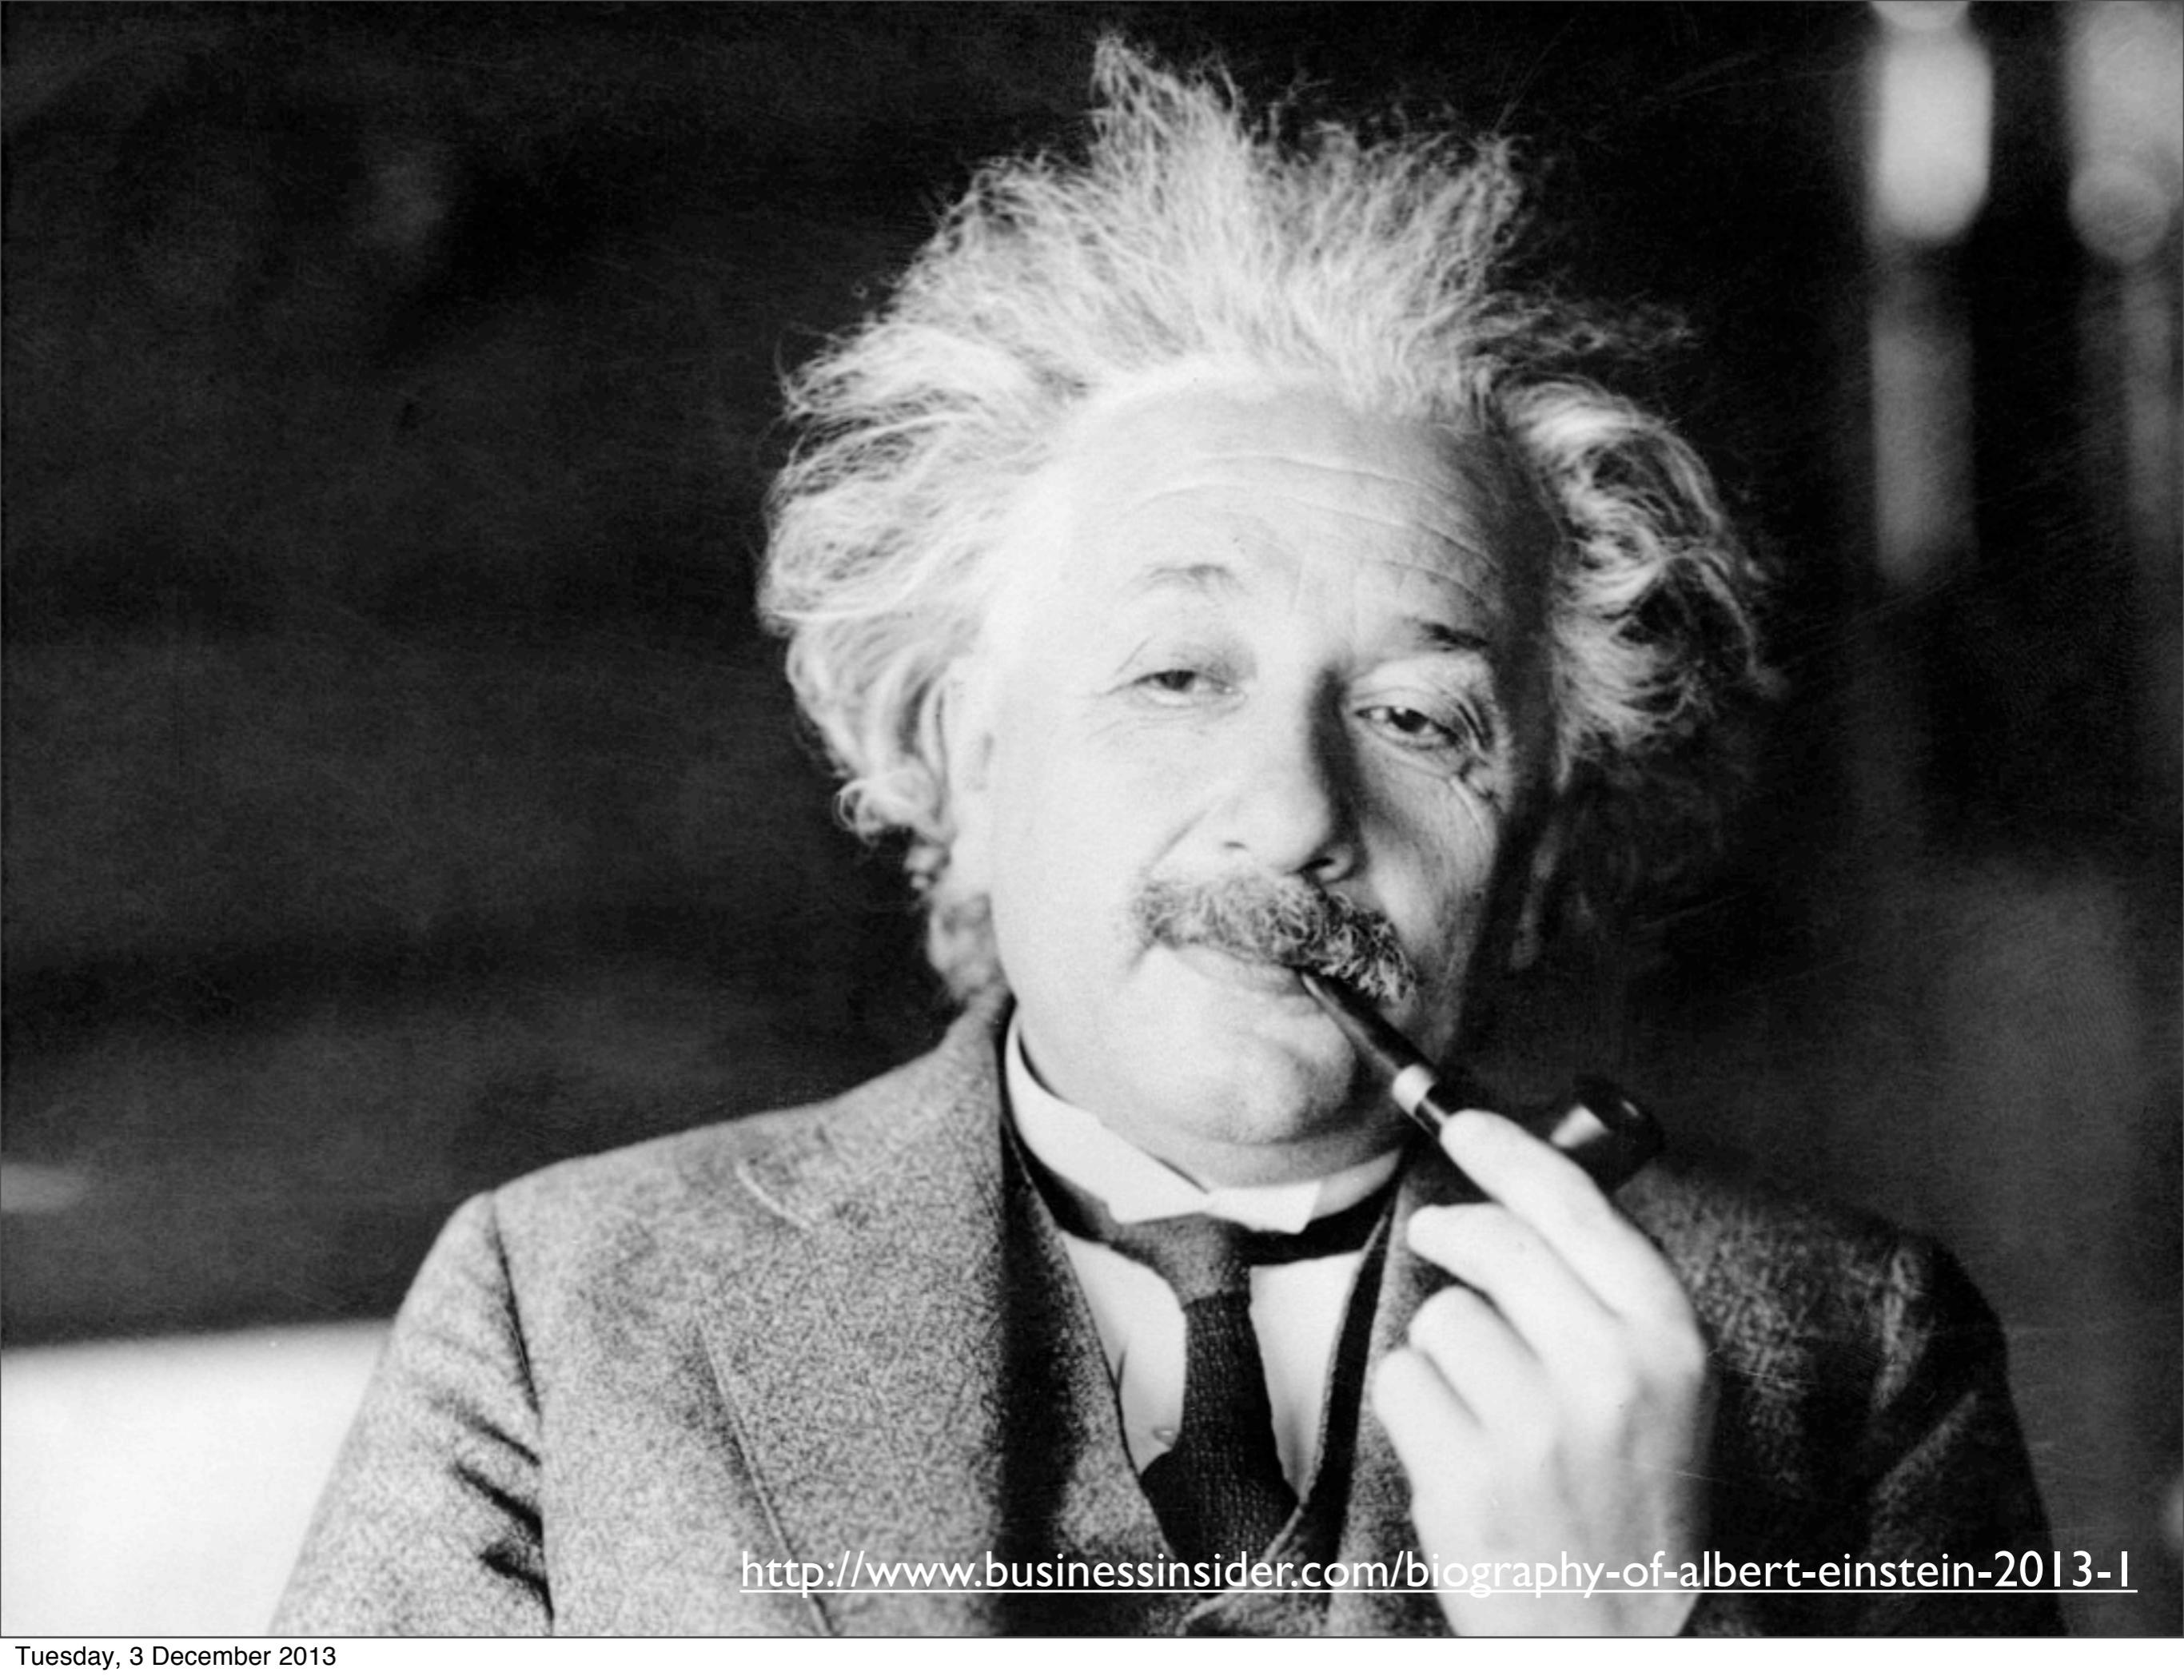
Solution #2 (Google)

Simply use atomic clocks to establish temporal ordering of events and distributed transactions ...



Solution #3 (Special Relativity)

Simply trade time for space!



<http://www.businessinsider.com/biography-of-albert-einstein-2013-1>

G-Counter

A counter that can
only GROW

G-Counter insight:
Store a separate
counter for each
machine

G-Counter

Machine A

A:0

B:0

Machine B

A:0

B:0

A machine can only
increment its own
counter

G-Counter

Machine A

A:5

B:0

Machine B

A:0

B:7

Merge is simply the
max

G-Counter

Machine A

A:5

B:7

Machine B

A:5

B:7

The counter's value
is simply the
total

G-Counter

Machine A

A:5

B:7

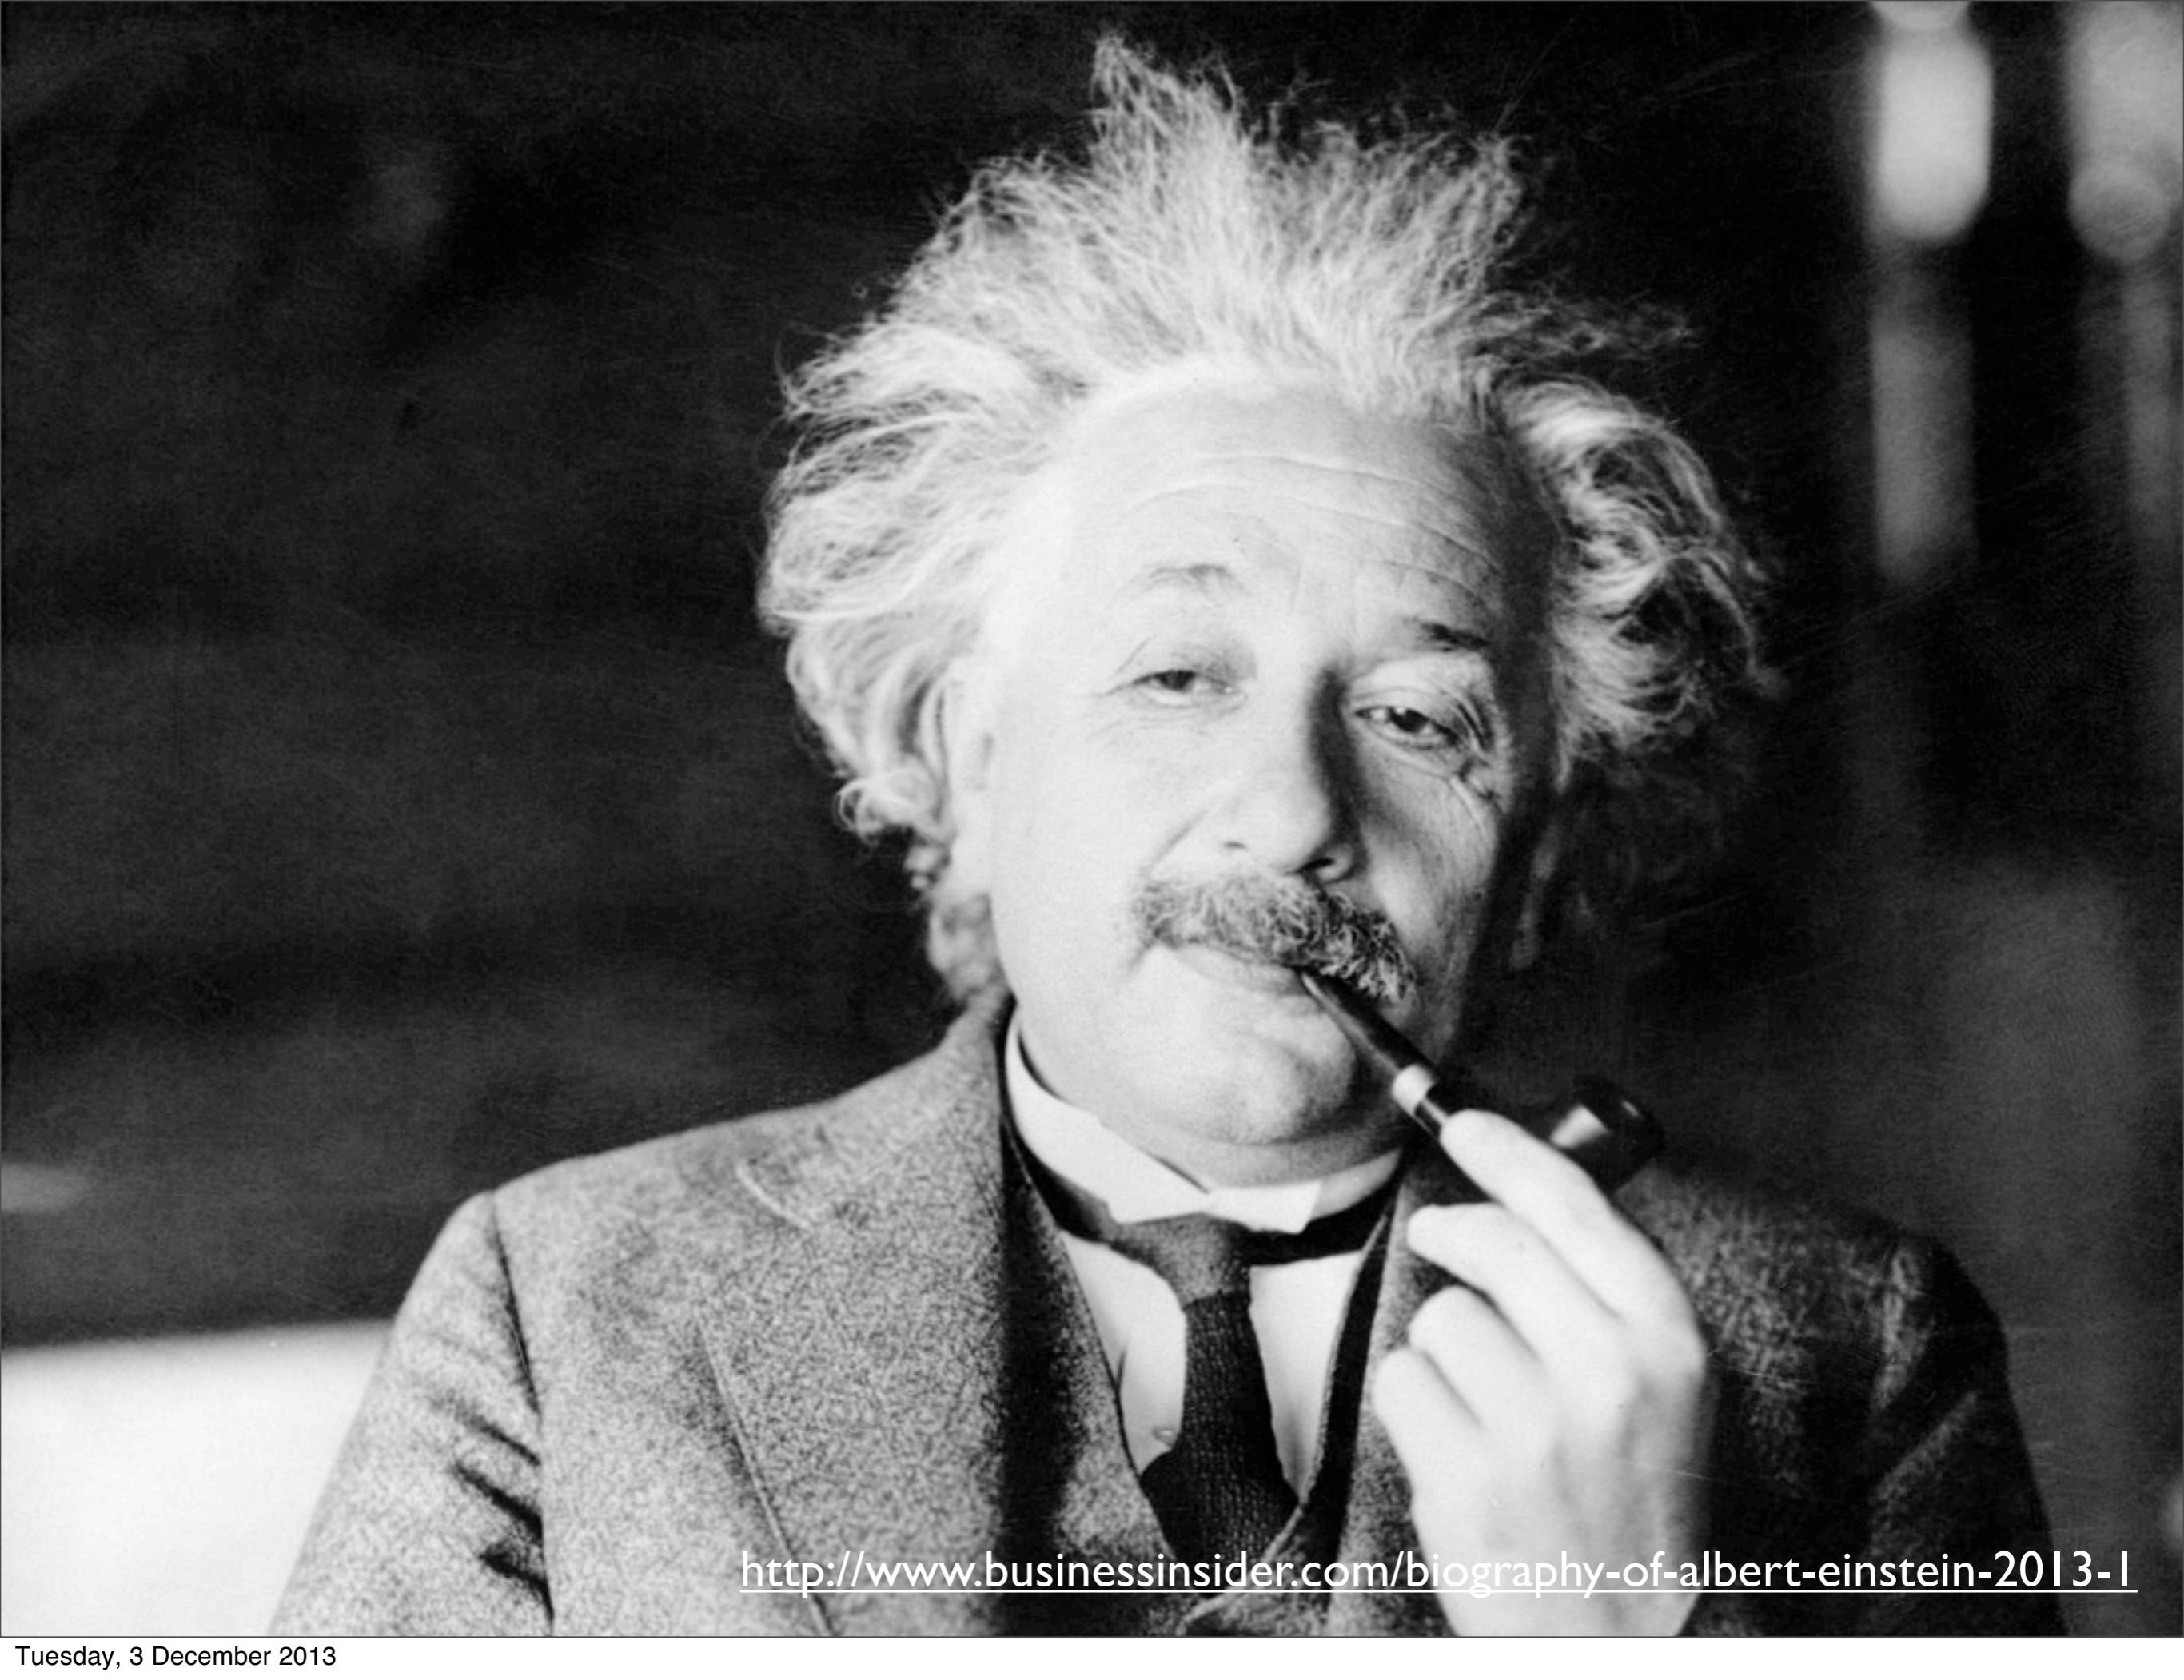
Machine B

Total A:5

is 10 B:7

We have a
distributed
eventually-consistent
increment-only
counter

We have used
space to become
invariant to time



<http://www.businessinsider.com/biography-of-albert-einstein-2013-1>

Can we abstract
this idea?

```
trait GCounter[Id, Elt] {  
  
  def inc(id: Id, amt: Elt)  
  
  def total: Elt  
  
  def merge(c: GCounter[Id,  
Elt]): GCounter[Id, Elt]  
}
```

total requires Elt has +

inc requires Elt has +,0

+ must be
**Invariant to
order**

Formally:

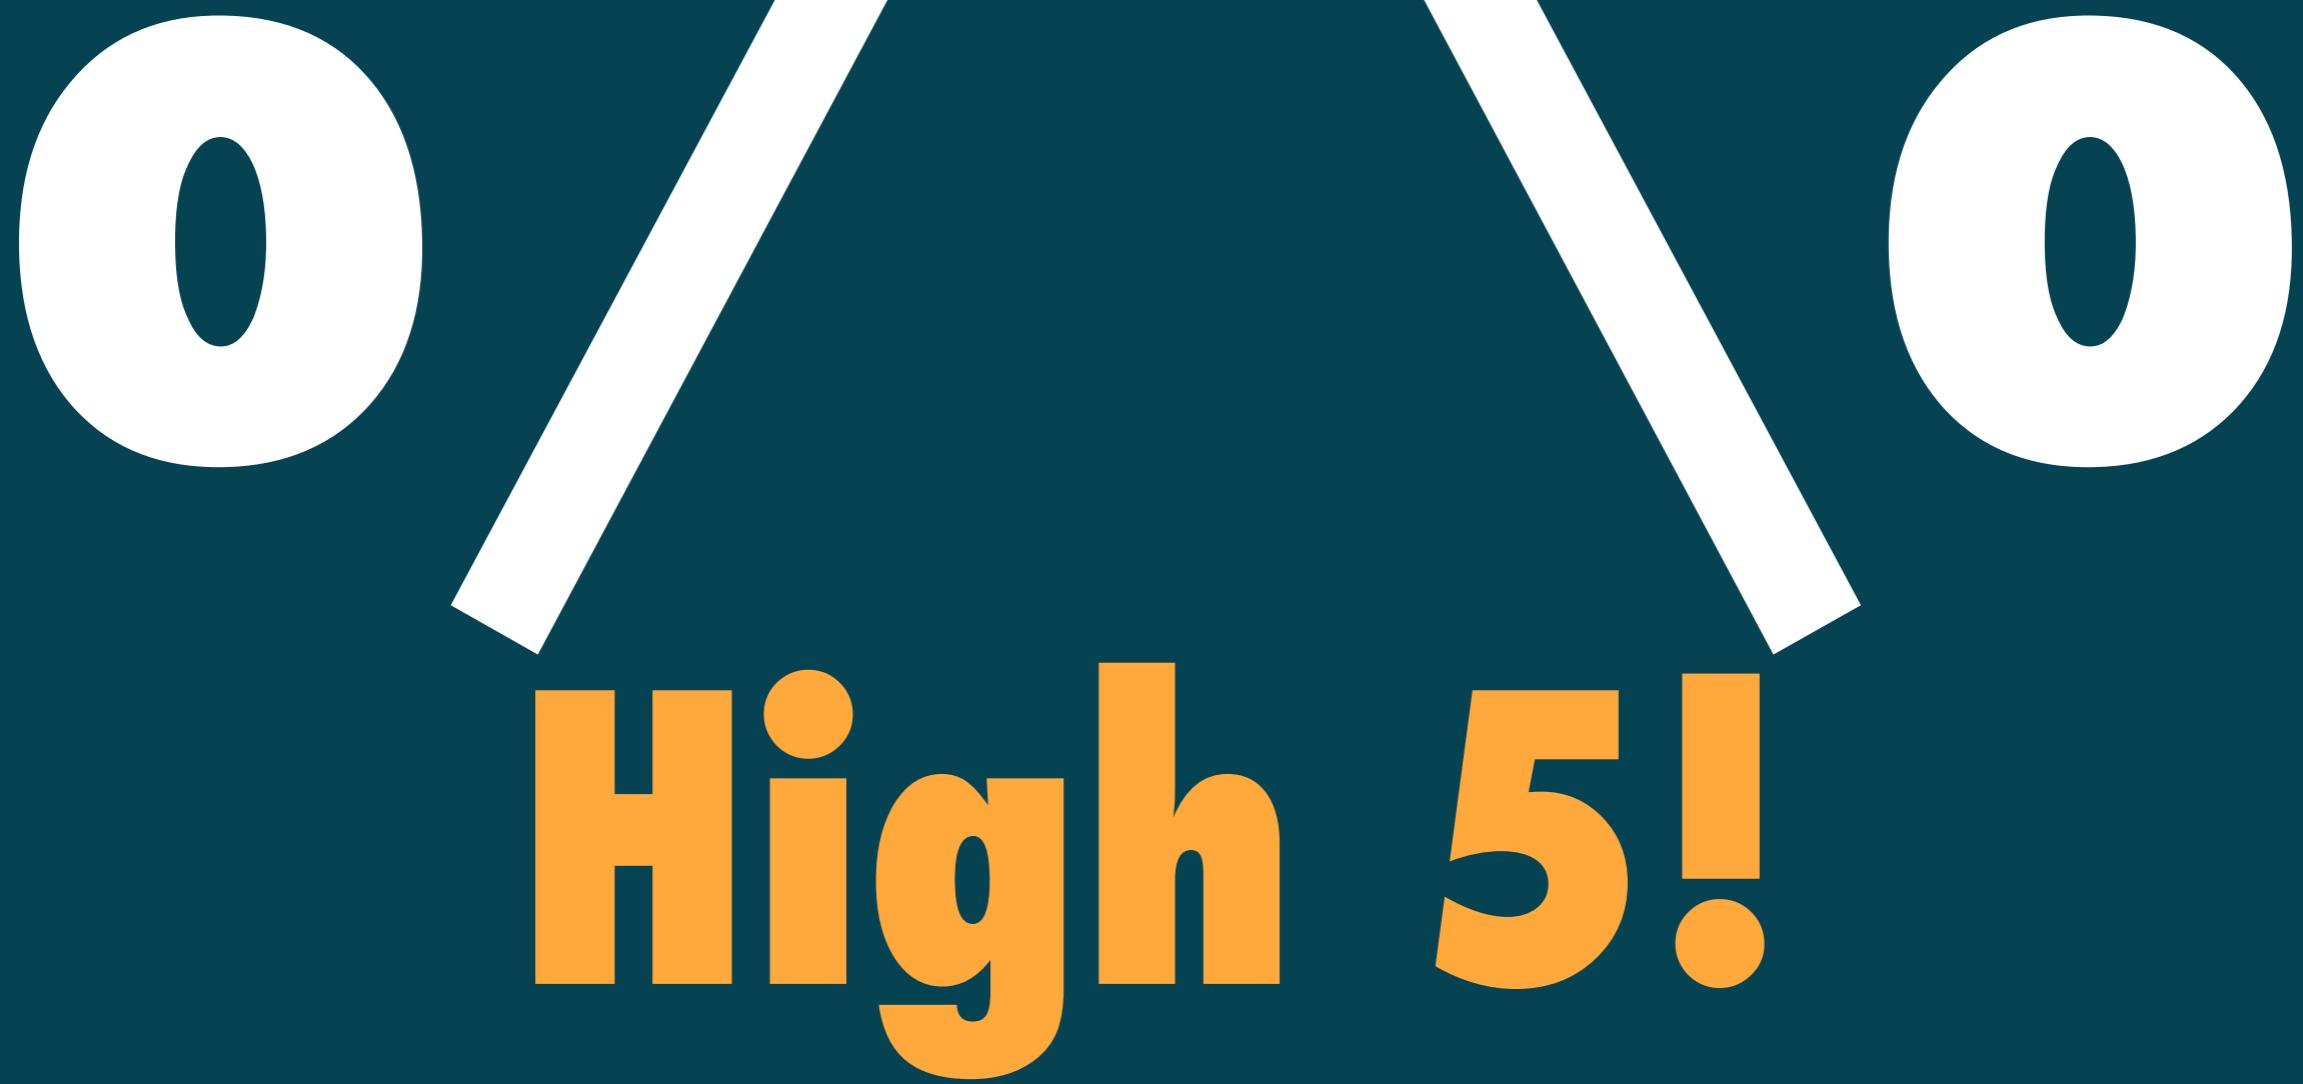
Associative

$$(x \bullet y) \bullet z = x \bullet (y \bullet z)$$

Formally:
Commutative

$$x \bullet y = y \bullet x$$

A Commutative Monoid!



A graphic featuring a dark teal background. In the center, the word "High 5!" is written in large, bold, orange letters. Above the text is a yellow asterisk (*). Two white diagonal bars extend from behind the asterisk, one pointing up and left, the other up and right. To the left of the first bar is a large white circle containing a dark teal oval. To the right of the second bar is another large white circle containing a dark teal oval.

High 5!

```
def inc(id: Id, amt: Elt)  
(implicit m: Monoid[Elt])
```

```
def total(implicit m:  
Monoid[Elt]): Elt
```

merge requires

Elf has max

**max must be
Invariant to
order**

Formally:

Associative

$$(x \bullet y) \bullet z = x \bullet (y \bullet z)$$

Formally:
Commutative

$$x \bullet y = y \bullet x$$

**max must
Converge to the
correct value**

Formally:
Idempotent

$$x \bullet x = x$$

An Idempotent
Commutative
Monoid!!!

```
def merge(c: GCounter[Id,  
Elt])(implicit m: Monoid[Elt  
@@ Max]): GCounter[Id, Elt]
```

Number (+)

Number (*)

Tuple

Map

Option

Average

Moving average

t-digest

Set (intersection)

Set (union)

Hyperloglog

Bloom filter

Count-min

Vector

Q-Tree

SGD

G-Set

Machine A

A:{x}

B:{}

Machine B

A:{}

B:{y,z}

G-Set Merge

Machine A

A:{x}
B:{y,z}

Machine B

A:{x}
B:{y,z}

G-Set Total

Machine A

A:{x} Set is
B:{y, z} {x, y, z} B:{y, z}

Machine B

PN-Counter

**A counter that can
GROW and
SHRINK**

**Can't use a G-
Counter as we can't
use max to merge**

Use
Two
G-counters!

PN-Counter

Machine A
Additions

A: 4, B: 2

Subtractions

A: 5, B: 3

Machine B
Additions

A: 4, B: 7

Subtractions

A: 3, B: 4

Merge is simply the
MAX

PN-Counter

Machine A
Additions

A: 4, B: 7

Subtractions

A: 5, B: 4

Machine B
Additions

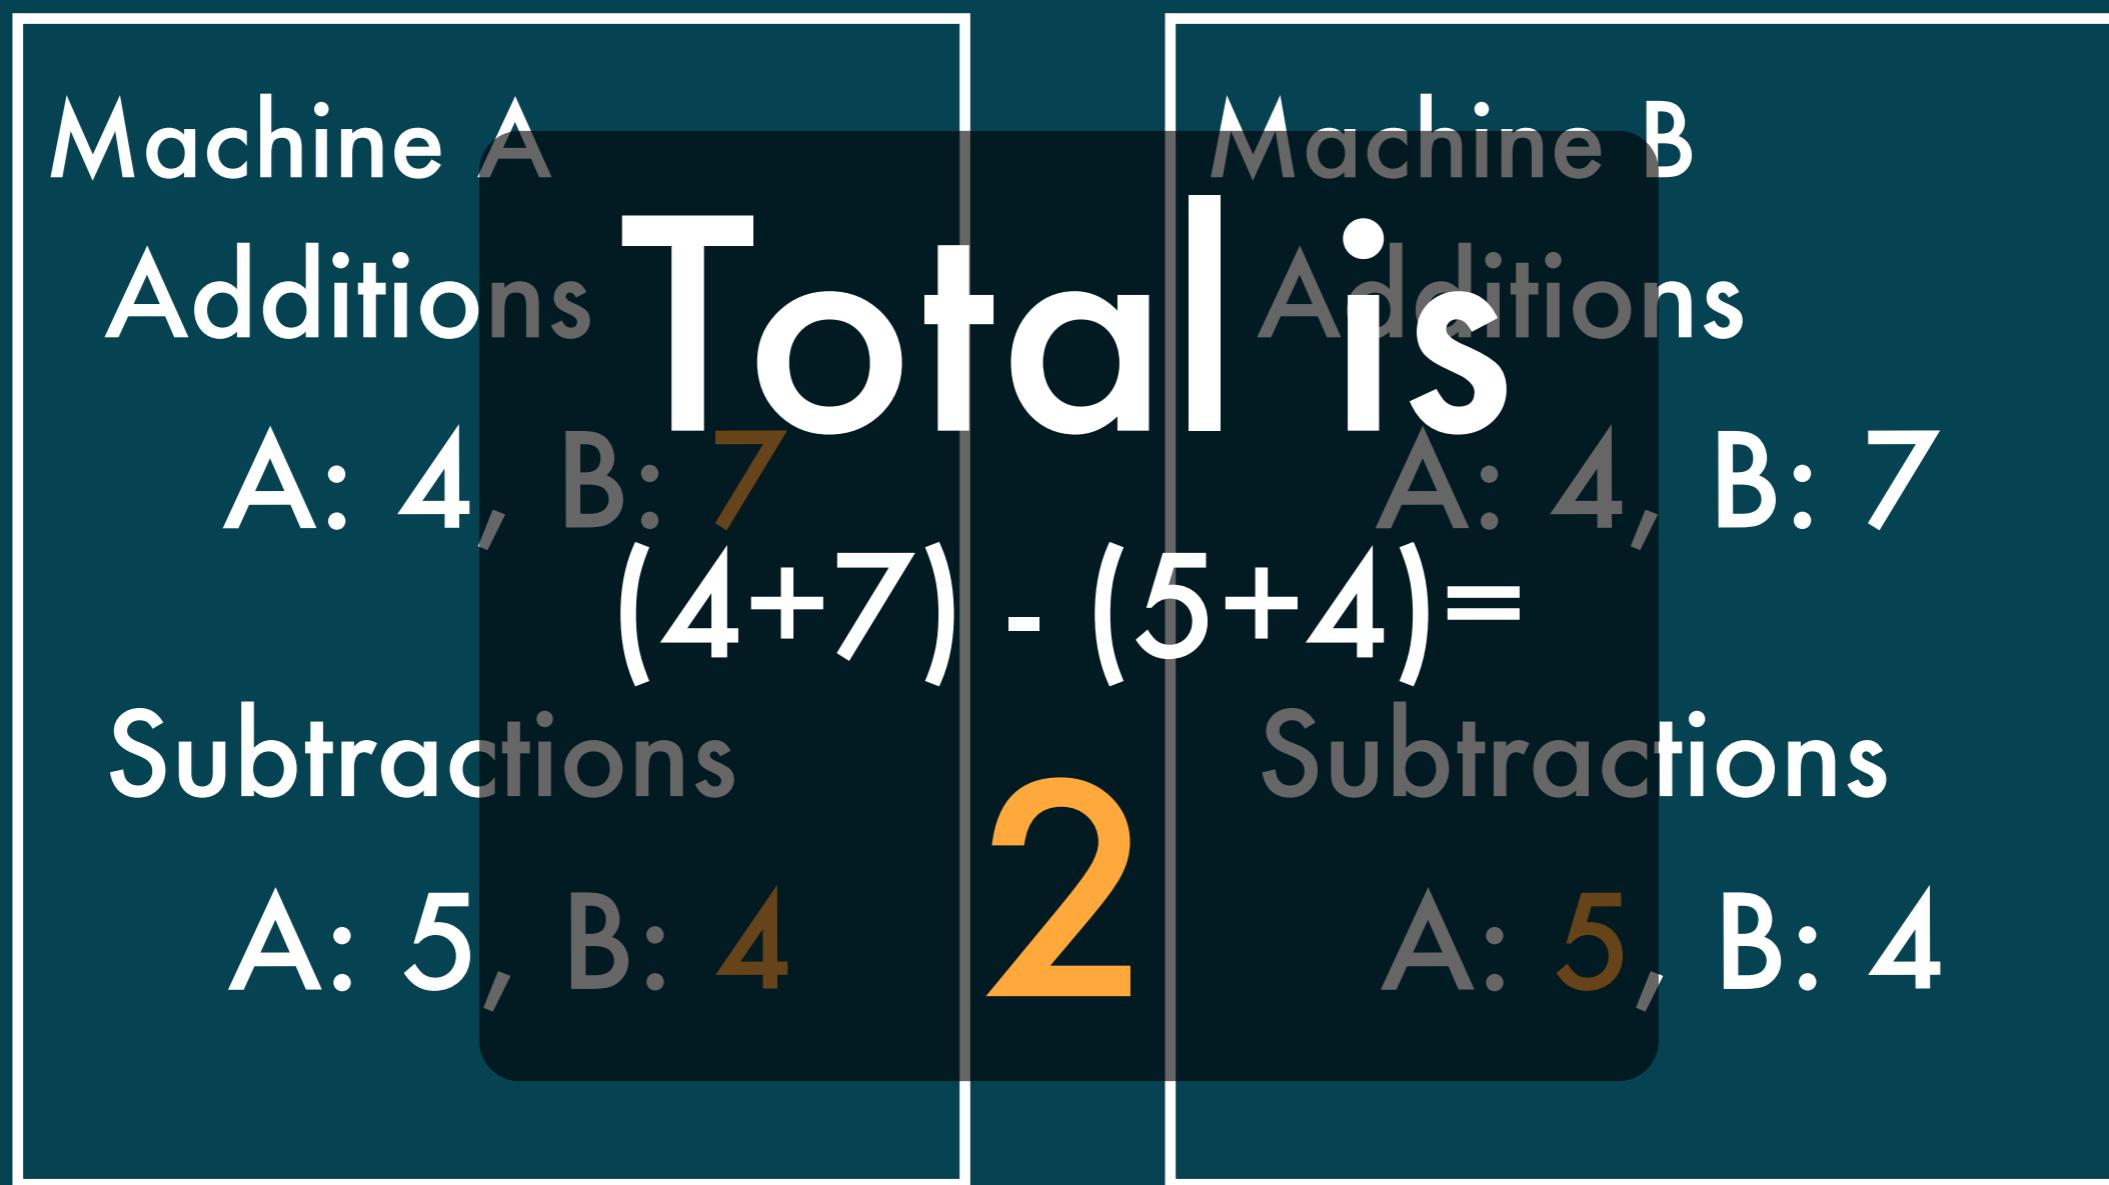
A: 4, B: 7

Subtractions

A: 5, B: 4

The counter's value
is simply the
TOTAL

PN-Counter



```
trait PNCounter[Id, Elt] {  
  
  def inc(id: Id, amt: Elt)  
  
  def dec(id: Id, amt: Elt)  
  
  def total: Elt  
  
  def merge(c: GCounter[Id, Elt]):  
    GCounter[Id, Elt]  
}
```

**PN-Counter requires
Elt has addition,
zero, and
subtraction**

A Commutative
Group!

```
trait PNCounter[Id,Elt] {  
  
  def inc(id: Id, amt: Elt)(implicit m: Monoid[Elt])  
  
  def dec(id: Id, amt: Elt)(implicit m: Monoid[Elt])  
  
  def total(implicit m: Group[Elt]): Elt  
  
  def merge(c: GCounter[Id, Elt])(implicit m: Monoid[Elt @@ Max]): GCounter[Id, Elt]  
}
```

**Numbers are clearly
a commutative
group**

**Sets with union and
set difference are a
commutative group**

PN-set

2P-Set

Machine A
Additions

A: {x}, B: {y}

Subtractions

A: {x}, B: {}

Machine B
Additions

A: {x}, B: {y, z}

Subtractions

A: {}, B: {y}

2P-Set Merge

Machine A
Additions

A: {x}, B: {y, z}

Subtractions

A: {x}, B: {y}

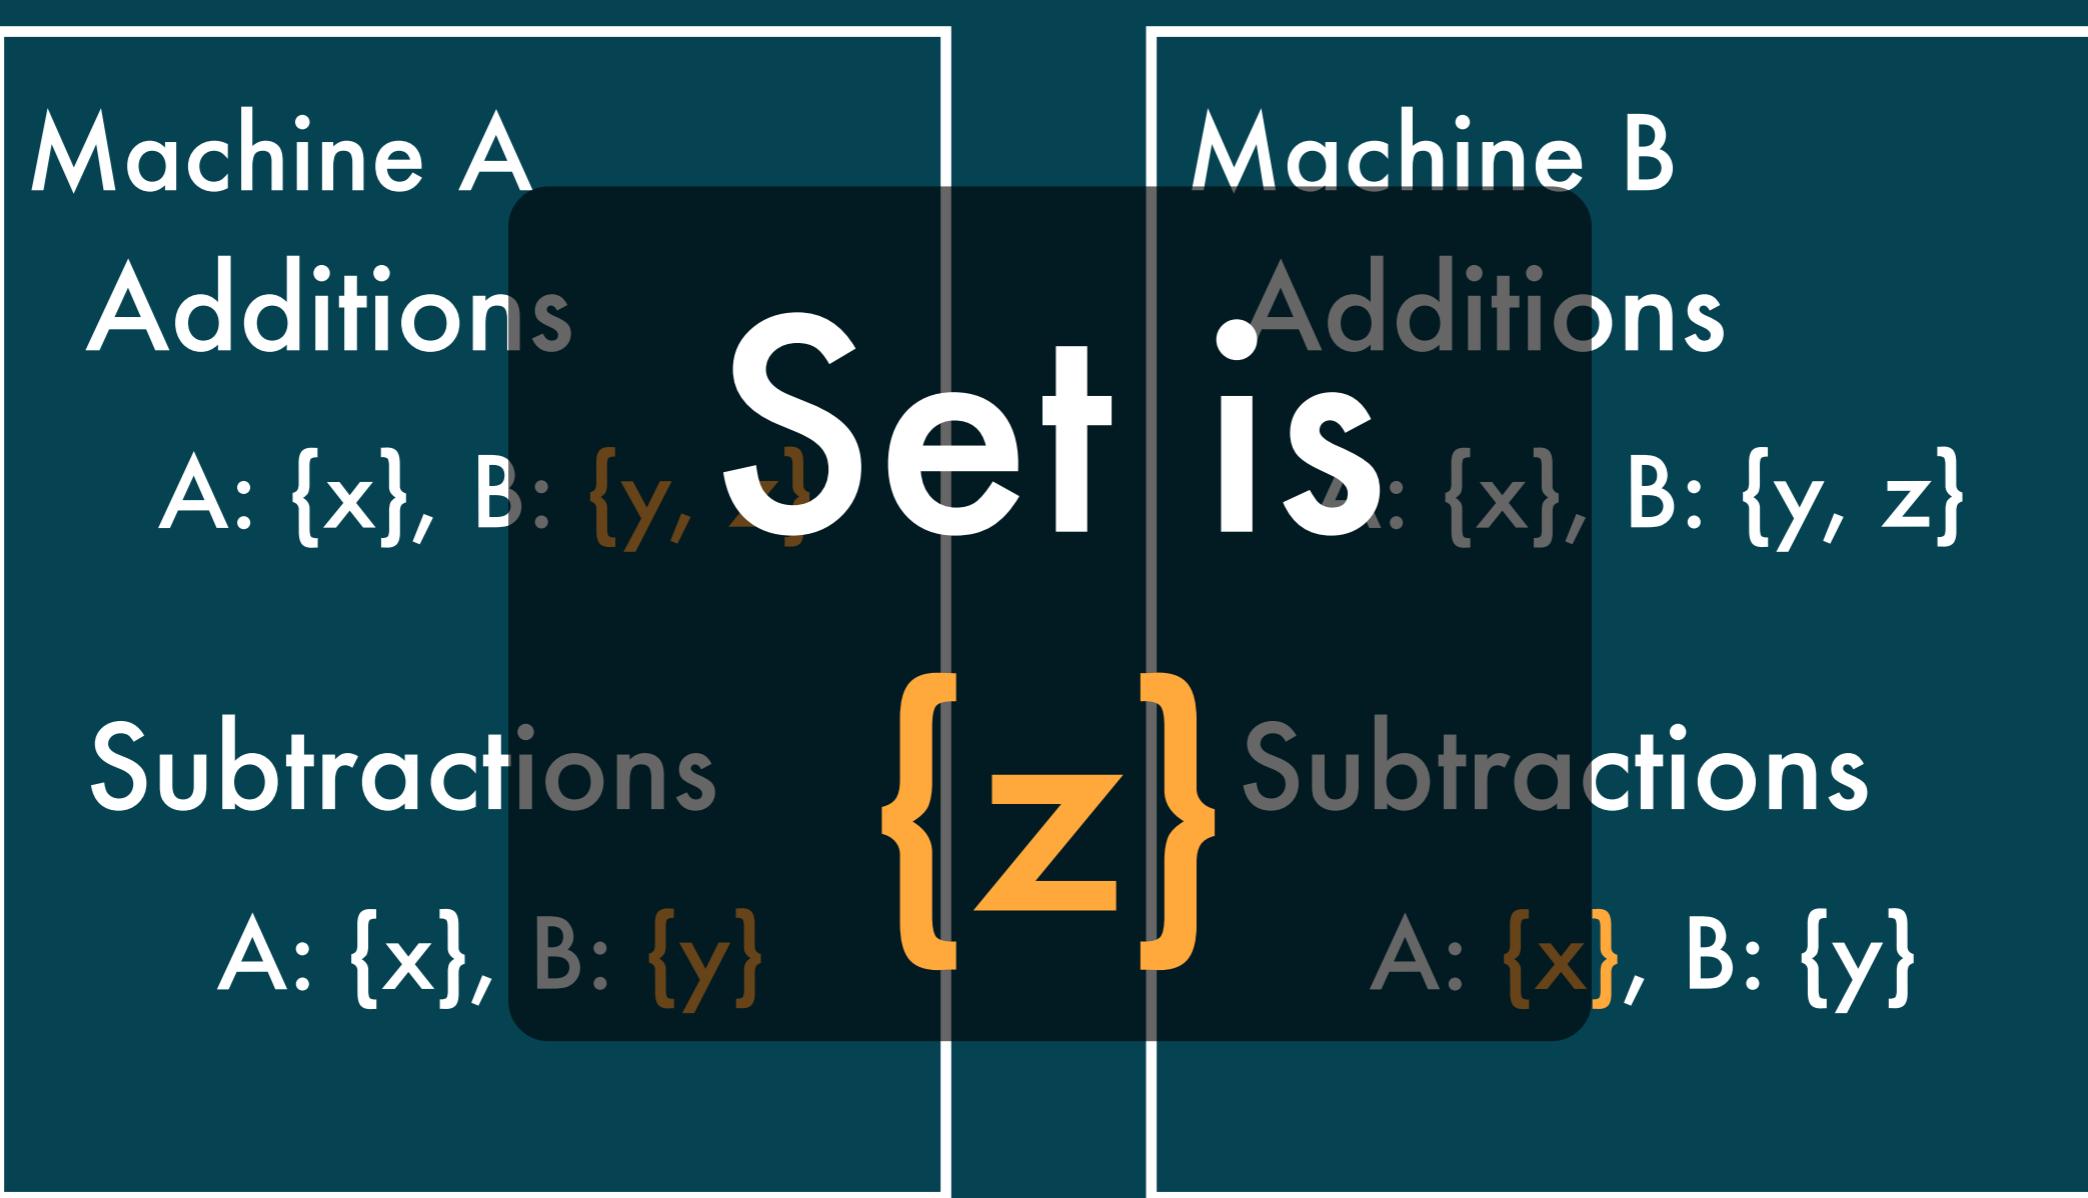
Machine B
Additions

A: {x}, B: {y, z}

Subtractions

A: {x}, B: {y}

2P-Set Total



**Deleted elements
stored indefinitely.**

Called

tombstones

**2P-Set allows
elements to be
added and removed
once**

C-Set

Store element and count

C-Set

Machine A

Additions

A: $\{(x, 2)\}$,
B: $\{(y, 1)\}$

Subtractions

A: $\{(x, 1)\}$,
B: $\{\}$

Machine B

Additions

A: $\{(x, 1)\}$,
B: $\{(y, 1), (z, 2)\}$

Subtractions

A: $\{\}$,
B: $\{(y, 1)\}$

C-Set Merge

Machine A

Additions

A: $\{(x, 2)\}$,
B: $\{(y, 1), (z, 2)\}$

Subtractions

A: $\{(x, 1)\}$,
B: $\{(y, 1)\}$

Machine B

Additions

A: $\{(x, 2)\}$,
B: $\{(y, 1), (z, 2)\}$

Subtractions

A: $\{(x, 1)\}$,
B: $\{(y, 1)\}$

C-Set Total

Machine A

Additions

A: $\{(x, 2)\}$,
B: $\{(y, 1), (z, 2)\}$

Subtractions

A: $\{(x, 1)\}$,
B: $\{(y, 1)\}$

Set

{x,
z}

Machine B

Additions

A: $\{(x, 2)\}$,
B: $\{(y, 1), (z, 2)\}$

Subtractions

A: $\{(x, 1)\}$,
B: $\{(y, 1)\}$

**C-Set allows
elements to be
added and removed
many times**

C-Set allows elements
to be removed
more times than
they have been added

OR-Set

**Store element and
unique token**

OR-Set

Machine A

Additions

A: $\{(x, \#a), (x, \#d)\},$
B: $\{(y, \#b)\}$

Subtractions

A: $\{(x, \#a)\},$
B: $\{\}$

Machine B

Additions

A: $\{(x, \#a)\},$
B: $\{(y, \#b), (z, \#c)\}$

Subtractions

A: $\{\},$
B: $\{(y, \#b)\}$

OR-Set Merge

Machine A

Additions

A: $\{(x, \#a), (x, \#d)\}$,
B: $\{(y, \#b), (z, \#c)\}$

Subtractions

A: $\{(x, \#a)\}$,
B: $\{(y, \#b)\}$

Machine B

Additions

A: $\{(x, \#a), (x, \#d)\}$,
B: $\{(y, \#b), (z, \#c)\}$

Subtractions

A: $\{(x, \#a)\}$,
B: $\{(y, \#b)\}$

OR-Set Total

Machine A

Additions

A: $\{(x, \#a), (x, \#d)\}$
B: $\{(y, \#b), (z, \#c)\}$

Subtractions

A: $\{(x, \#a)\},$
B: $\{(y, \#b)\}$

Set

Machine B

Additions

A: $\{(x, \#a), (x, \#d)\},$
B: $\{(y, \#b), (z, \#c)\}$

Subtractions
 $\{x, z\}$

A: $\{(x, \#a)\},$
B: $\{(y, \#b)\}$

**OR-Set works the
way we expect**

**From sets, build
trees, graphs,
etc.**

CRDTS vs The Real World

Strong Consistency Memory Usage Code

Strong Consistency

**Don't build your
billing platform on
CRDTs**

Memory Usage Tombstones Machine IDs

**Tombstones: Establish
causal order
and delete
(Bieniusa et al. 2012)**

**Tombstones: Prune
with heuristics
(often based on
time)**

Machine IDs: OR-Sets don't need them

**Machine IDs: Hierarchical
organisation allows
pruning
(Almeida & Baquero.
2013)**

**Code
Riak 2.0**

**Various open source
libraries**



Thank you!
Now go forth and
DISTRIBUTE!

<http://stjost.deviantart.com/art/Stomping-Off-Into-the-Sunset-277086274>

More:
noelwelsh.com